

# Per Entity Training Pipelines in Apache Beam

Jasper Van den Bossche  
ML6



We are a group of AI and machine learning experts building custom AI solutions.

Amongst our engineers we have several Apache Beam contributors.



# Agenda








- Development of ML applications
  - What is training?
  - What is MLOps?
- What does per entity training mean?
  - Training multiple models rather than a single model?
  - Why use a per entity strategy
- Example per entity training pipeline
- Bonus: Using trained models in a RunInference pipeline



What is machine learning model training?


# What is machine learning model training?

```
def contains_firefly():
    ...
```




Writing logic to detect the Beam mascot is almost impossible

# What is training a machine learning model?



# What is training a machine learning model?






How are machine learning  
applications built and deployed?




# MLOps: Level 0






# MLOps: Level 1






# MLOps: Level 2










What is per entity training?


## Example: Building multilingual chatbot





# What is per entity training?



## Example: Detect production defects using sensor data



## Example: Detect production defects using sensor data









# Why use a per entity strategy?




# Reduce Model Infrastructure Requirements



CPU Machine






Lightweight GPU



GPU Cluster




# Faster training & inference






# Address fairness and bias





# Easier to detect problems




.25	.14	.36	.25
.35	.45	.08	.12
.12	.23	.33	.32
.28	.18	.23	.31


Confusion Matrix




# Simpler models have the following advantages




Less powerful hardware required



Easier to address bias



Faster training & inference




Easier debugging



But there is one big problem:  
*How do I manage the training of  
all of these models?*



# Manage training pipelines





# The solution? Apache Beam!



- Apache Beam can handle *streaming* and *batch data*
- Apache Beam can easily *prepare data* for training
- Apache Beam can run on different *runners* depending on the model's *requirements*
- *Abstraction* in ML libraries allows us to train models with few lines of code



Let's look at an example of a per entity training pipeline



# Predicting incomes per education level



Age	Workclass	Education	Marital Status	Occupation	Relationship	Race	Sex	Hours per Week	Native Country	Compensation
25	Private	11th	Never-married	Machine-op-inspct	Own-child	Black	Male	40	USA	<=50K.
38	Private	HS-grad	Married-civ-spouse	Farming-fishing	Husband	White	Male	50	USA	<=50K.
28	Local-gov	Assoc-acdm	Married-civ-spouse	Protective-serv	Husband	White	Male	40	USA	>50K.
44	Private	Some-college	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	40	USA	>50K.
18	?	Some-college	Never-married	?	Own-child	White	Female	30	USA	<=50K.



# Pipeline overview






# Split data per education level



Age	...	Occupation	Education
29y	...	Accountant	Bachelor
31y	...	Engineer	Master
54y	...	Plumber	Bachelor
37y	...	Server	High School
47y	...	Barista	High School
22y	...	Cashier	Bachelor







# Train model per dataset



29y	...	Accountant	Bachelor
54y	...	Plumber	Bachelor
22y	...	Cashier	Bachelor
...	...	...	...

31y	...	Engineer	Master
...	...	...	...
...	...	...	...

37y	...	Server	High School
47y	...	Barista	High School
...	...	...	...





# Pipeline overview



```
with beam.Pipeline(options=pipeline_options) as pipeline:  
    _ = (  
        pipeline | "Read Data" >> beam.io.ReadFromText(known_args.input)  
        | "Split data to make List" >> beam.Map(lambda x: x.split(','))  
        | "Filter rows" >> beam.Filter(custom_filter)  
        | "Create Key" >> beam.ParDo(CreateKey())  
        | "Group by education" >> beam.GroupByKey()  
        | "Prepare Data" >> beam.ParDo(PrepareDataforTraining())  
        | "Train Model" >> beam.ParDo(TrainModel())  
        | "Save" >> fileio.WriteToFiles(path=known_args.output,  
                                         sink=ModelSink()))
```



# Step 1: Data preparation



```
def custom_filter(element):
    return len(element) == 15 and '?' not in element \
        and ' Bachelors' in element or ' Masters' in element \
        or ' Doctorate' in element
```



# Step 1: Data preparation



```
class PrepareDataforTraining(beam.DoFn):
    def process(self, element, *args, **kwargs):
        key, values = element

        #Convert to dataframe
        df = pd.DataFrame(values)
        last_ix = len(df.columns) - 1
        X, y = df.drop(last_ix, axis=1), df[last_ix]

        # select categorical and numerical features
        cat_ix = X.select_dtypes(include=['object', 'bool']).columns
        num_ix = X.select_dtypes(include=['int64', 'float64']).columns

        # label encode the target variable to have the classes 0 and 1
        y = LabelEncoder().fit_transform(y)

        yield (X, y, cat_ix, num_ix, key)
```



## Step 2: Training the models



```
class TrainModel(beam.DoFn):

    def process(self, element, *args, **kwargs):
        X, y, cat_ix, num_ix, key = element
        steps = [('c', OneHotEncoder(handle_unknown='ignore'), cat_ix),
                 ('n', MinMaxScaler(), num_ix)]

        # one hot encode categorical, normalize numerical
        ct = ColumnTransformer(steps)

        # wrap the model in a pipeline
        pipeline = Pipeline(steps=[('t', ct), ('m', DecisionTreeClassifier())])
        pipeline.fit(X, y)

        yield (key, pipeline)
```



## Step 3: Saving models




```
class ModelSink(fileio.FileSink):
    def open(self, fh):
        self._fh = fh

    def write(self, record):
        _, trained_model = record
        pickled_model = pickle.dumps(trained_model)
        self._fh.write(pickled_model)

    def flush(self):
        self._fh.flush()
```



# Extending the pipeline





# Extending pipeline with metrics




```
class EvaluateModel(beam.DoFn):
    def __init__(self, model_uri):
        file = FileSystems.open(model_uri, 'rb')
        self.model = pickle.load(file)

    def process(self, element, *args, **kwargs):
        inputs, labels = element
        predictions = self.model.predict(inputs)
        accuracy = sklearn.metrics.accuracy_score(y_pred=predictions,
                                                y_true=labels)
        f1 = sklearn.metrics.f1_score(y_pred=predictions, y_true=labels)
        recall = sklearn.metrics.recall_score(y_pred=predictions, y_true=labels)

        file = FileSystems.open(f'model_uri_metrics', 'wb')
        file.writelines([f'accuracy: {accuracy}', f'f1: {f1}', f'recall:
{recall}'])
```



# How does this pipeline fit in the MLOps architecture?






Let's try out our model using the  
RunInference transform



# Bonus: Inference in Apache Beam





# Summary



- Apache Beam is more and more becoming technology that can be used in advanced MLOps setups
- Per entity strategy has several advantages
  - Requires less powerful hardware
  - Faster training and inference
  - Easier to address bias
  - Easier to debug
- Apache Beam a perfect candidate for per entity training pipelines thanks to
  - Excellent for data preprocessing and preparation
  - Different runners depending on model requirements
  - Abstraction in ML libraries that make it easy to train a model

Jasper Van den Bossche

# QUESTIONS?

<https://www.linkedin.com/in/jasper-van-den-bossche/>

<https://github.com/jaxpr>

<https://www.ml6.eu/>