

BEAM
SUMMIT

Complex Event Processing with State and Timers

Israel Herraiz, Miren Esnaola
Google

- Have a Python development environment ready, with **Python 3.10**
- Use your favourite IDE (Visual Studio Code, PyCharm, ...)
- Clone the workshop repo

github.com/iht/beam-state-timers-quickstart

- Install dependencies with

```
pip install -r requirements.txt
```

- Run the tests (they fail until you write the solution 😊) with `pytest`.
- Have a look at the file `my_pipeline/pipeline.py` and fill the gaps



Agenda



- What is stateful stream processing?
- Uses cases
- State properties
- Types of state
- Example
- Timers

Stateful stream processing — What is it?

Stateful stream processing is a subset of stream processing in which the computation maintains contextual state. This state is used to store information derived from the previously-seen events.



Stateful stream processing — Use cases

Most non-trivial stream processing applications require stateful event processing:

Personalization

A video streaming service could use it to track a user's past viewing history and use this information to make recommendations for movies and TV shows.

Fraud detection

A financial institution could use it to track a user's past transactions and use this information to detect unusual activity that might indicate fraud.

Supply chain management

A logistics company could use it to track the location and status of packages in real-time, and to optimize the routing of packages based on past delivery times and delays.

State Properties

It is identified by the name that must be **unique** through the transform.

It must remain **local** to the transform.

It can contain **different types of objects**: scalar values, collections or maps.

It works **per key**.

It is **bound to a window**.

How it is stored **depends on the runner** implementation.
For the Direct Runner state is stored in memory.

Types of State



ReadModifyWriteState

A readable state cell containing a single value.

CombiningValueState

A readable state cell defined by a function, accepting multiple input values, combining them and producing a single output value.

BagState

A readable state cell containing a bag of values. Items can be added to the bag and the contents read out.

Types of State — ReadModifyWriteState



```
class ReadModifyWriteStateDoFn(DoFn):  
    STATE_SPEC = ReadModifyWriteStateSpec('num_elements', VarIntCoder())  
  
    def process(self, element, state=DoFn.StateParam(STATE_SPEC)):  
        # Read the number element seen so far for this user key.  
        current_value = state.read() or 0  
        state.write(current_value + 1)  
  
_ = (p | 'Read per user' >> ReadPerUser()  
     | 'state pardo' >> beam.ParDo(ReadModifyWriteStateDoFn()))
```

Types of State — CombineValueState



```
class CombiningStateDoFn(DoFn):  
    SUM_TOTAL = CombiningValueStateSpec('total', sum)  
    def process(self, element, state=DoFn.StateParam(SUM_TOTAL)):  
        state.add(1)  
  
_ = (p | 'Read per user' >> ReadPerUser()  
    | 'Combine state pardo' >>  
beam.ParDo(CombiningStateDoFn()))
```

Types of State — BagState



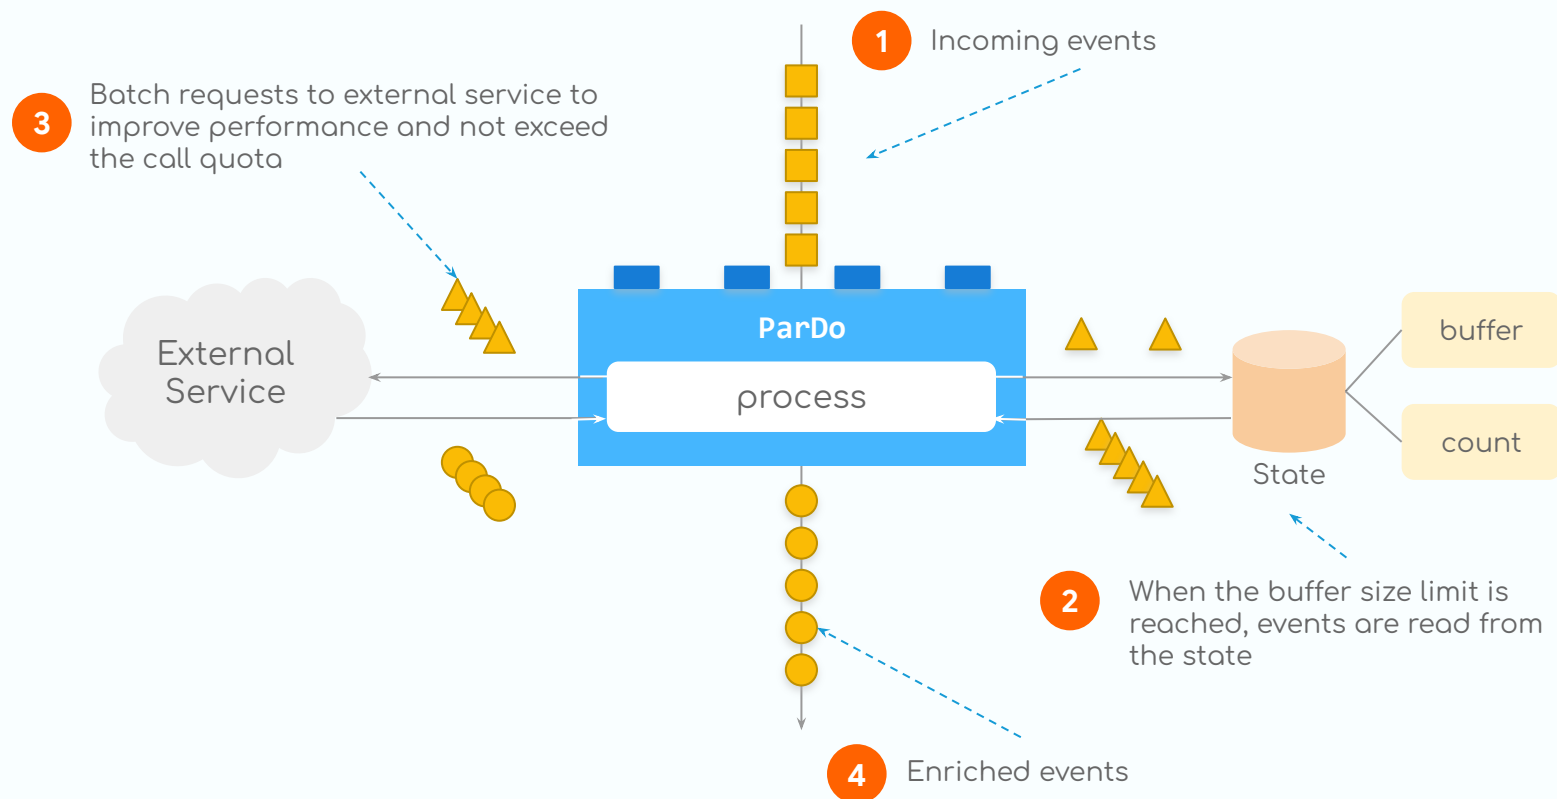
```
class BagStateDoFn(DoFn):
    ALL_ELEMENTS = BagStateSpec('buffer', coders.VarIntCoder())
    def process(self, element_pair, state=DoFn.StateParam(ALL_ELEMENTS)):
        state.add(element_pair[1])
        if should_fetch():
            all_elements = list(state.read())
            process_values(all_elements)
            state.clear()
_ = (p | 'Read per user' >> ReadPerUser()
    | 'Bag state pardo' >> beam.ParDo(BagStateDoFn()))
```



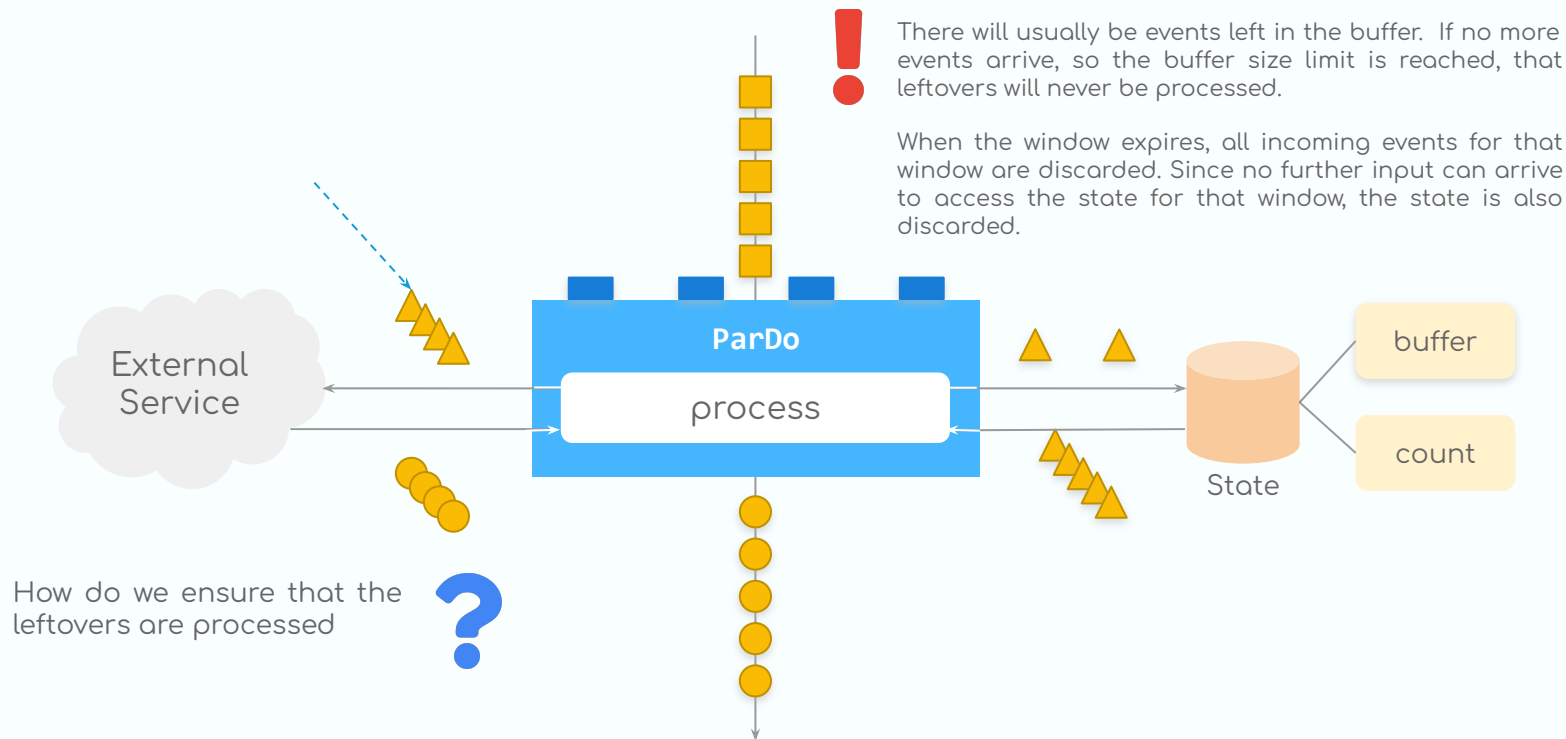
Umm, let's look at a stateful event processing example...

Suppose you have a pipeline with a ton of data coming in and you need to enrich each event by with the information returned by a external system. You can't just issue a call per event, performance would be terrible and it would blow your quota with that external system. What shall you do?

Example



Example






```
class StatefulBufferingFn(beam.DoFn):
    MAX_BUFFER_SIZE = 500;
    BUFFER_STATE = BagStateSpec('buffer', EventCoder())
    COUNT_STATE = CombiningValueStateSpec('count',
                                           VarIntCoder(),
                                           combiners.SumCombineFn())

    def process(self, element,
                buffer_state=beam.DoFn.StateParam(BUFFER_STATE),
                count_state=beam.DoFn.StateParam(COUNT_STATE)):
        buffer_state.add(element)
        count_state.add(1)
        count = count_state.read()

        if count >= MAX_BUFFER_SIZE:
            for event in buffer_state.read():
                yield event
            count_state.clear()
            buffer_state.clear()
```

Event-time Timers

Callback when the watermark reaches some threshold.



Processing-time Timers

Callback after a certain amount of time has elapsed.

Timers — Event-time Timer



```
class EventTimerDoFn(DoFn):  
    ALL_ELEMENTS = BagStateSpec('buffer', coders.VarIntCoder())  
    TIMER = TimerSpec('timer', TimeDomain.WATERMARK)  
  
    def process(self,  
                element_pair,  
                t = DoFn.TimestampParam,  
                buffer = DoFn.StateParam(ALL_ELEMENTS),  
                timer = DoFn.TimerParam(TIMER)):  
        buffer.add(element_pair[1])  
        # Set an event-time timer to the element timestamp.  
        timer.set(t)  
  
    @on_timer(TIMER)  
    def expiry_callback(self, buffer = DoFn.StateParam(ALL_ELEMENTS)):  
        state.clear()  
  
_ = (p | 'Read per user' >> ReadPerUser()  
    | 'EventTime timer pardo' >> beam.ParDo(EventTimerDoFn()))
```

Timers — Processing Timer



```
class ProcessingTimerDoFn(DoFn):

    ALL_ELEMENTS = BagStateSpec('buffer', coders.VarIntCoder())
    TIMER = TimerSpec('timer', TimeDomain.REAL_TIME)

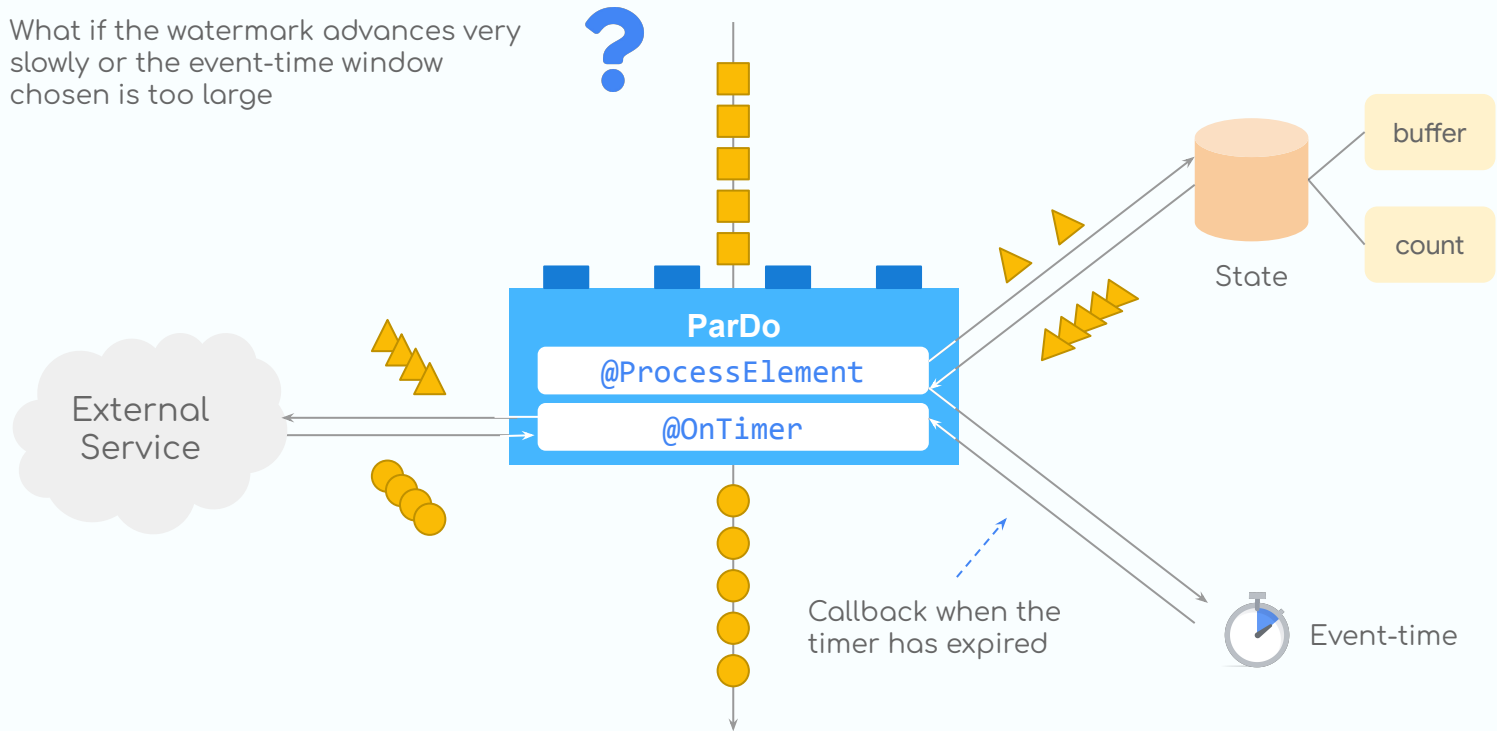
    def process(self,
                element_pair,
                t = DoFn.TimestampParam,
                buffer = DoFn.StateParam(ALL_ELEMENTS),
                timer = DoFn.TimerParam(TIMER)):
        buffer.add(element_pair[1])
        # Set a timer to go off 30 seconds in the future.
        timer.set(Timestamp.now() + Duration(seconds=30))

    @on_timer(TIMER)
    def expiry_callback(self, buffer = DoFn.StateParam(ALL_ELEMENTS)):
        state.clear()

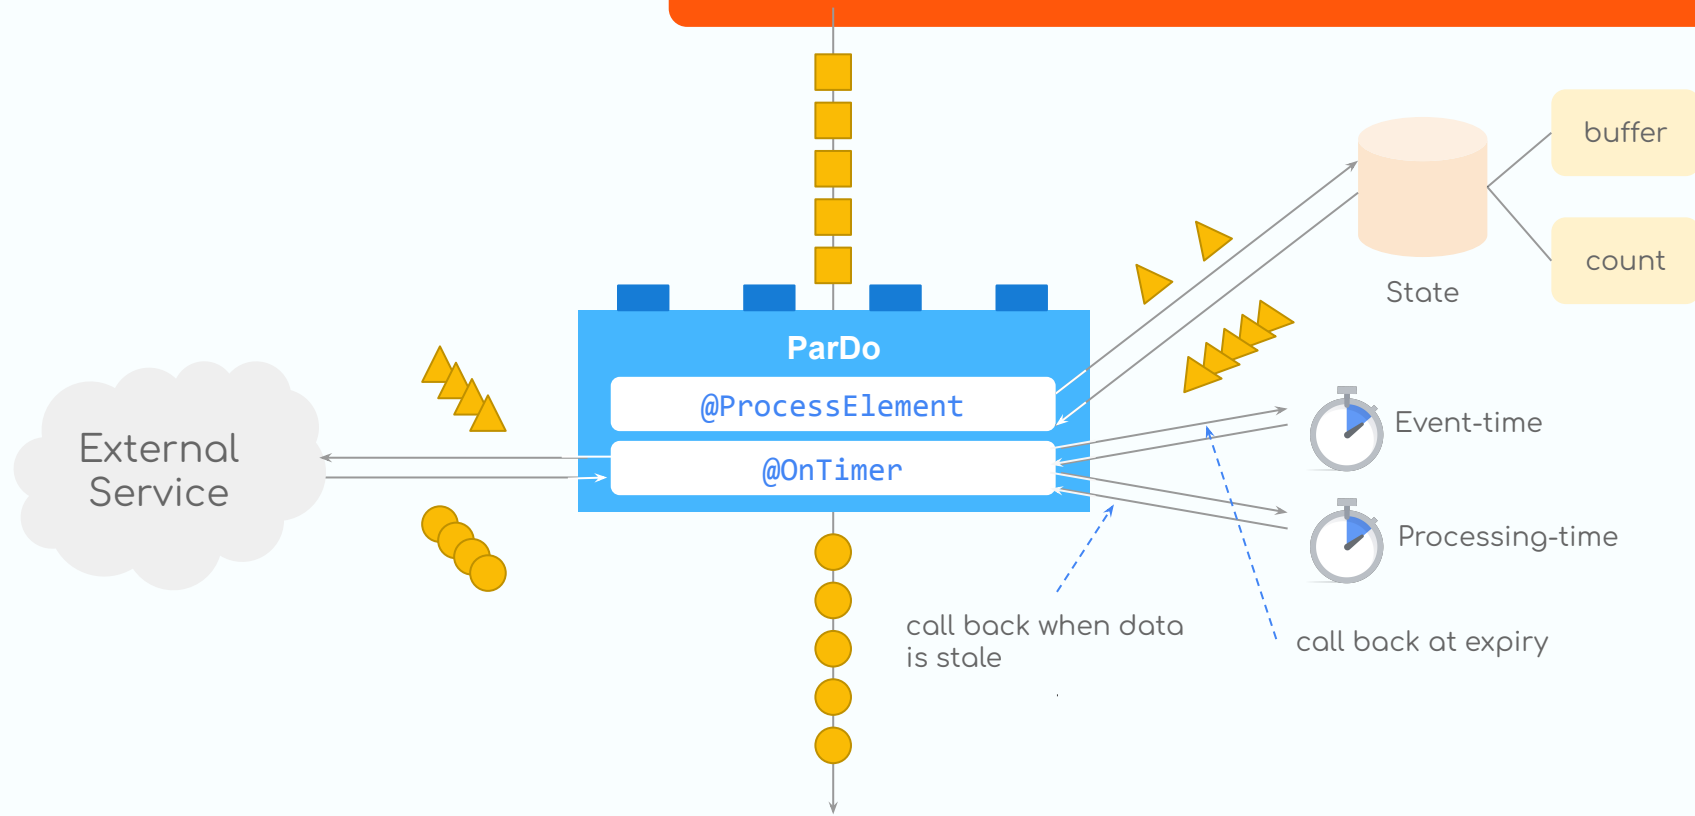
_ = (p | 'Read per user' >> ReadPerUser()
    | 'EventTime timer pardo' >> beam.ParDo(ProcessingTimerDoFn()))
```

Timers

What if the watermark advances very slowly or the event-time window chosen is too large



Timers





```
class StatefulBufferingFn(beam.DoFn):
    ...
    STALE_TIMER = TimerSpec(stale, TimeDomain.REAL_TIME)

    def process(self, element,
                w=beam.DoFn.WindowParam,
                buffer_state=beam.DoFn.StateParam(BUFFER_STATE),
                count_state=beam.DoFn.StateParam(COUNT_STATE),
                expiry_timer=beam.DoFn.TimerParam(EXPIRY_TIMER),
                stale_timer=beam.DoFn.TimerParam(STALE_TIMER)):
        if count_state.read() == 0:
            # We set an absolute timestamp here (not an offset like in the Java SDK)
            stale_timer.set(time.time() + StatefulBufferingFn.MAX_BUFFER_DURATION)
            ... same logic as above ...
```

```
@on_timer(STALE_TIMER)
def stale(self,
          buffer_state=beam.DoFn.StateParam(BUFFER_STATE),
          count_state=beam.DoFn.StateParam(COUNT_STATE)):
    events = buffer_state.read()

    for event in events:
        yield event

    buffer_state.clear()
    count_state.clear()
```

added an additional processing-time timer in case the buffer is filling too slowly



It's time for a lab, folks!!!



Input messages

```
{  
  "latitude" : 40.77405,  
  "longitude" : -73.9638,  
  "meter_increment" : 0.024726477,  
  "meter_reading" : 6.428884,  
  "passenger_count" : 5,  
  "point_idx" : 260,  
  "ride_id" : "ccf021d0-ec37-41f1-9637-cf8bcfbcbb2d",  
  "ride_status" : "enroute",  
  "timestamp" : "2023-06-11T09:40:03.15611-04:00"  
}
```


Output messages

```
{  
  "ride_id" : "ccf021d0-ec37-41f1-9637-cf8bcfbcbb2d",  
  "start_time" : "2023-06-11T09:40:03.15611-04:00"  
  "end_time" : "2023-06-11T15:47:03.16611-04:00"  
  "start_status" : "pickup"  
  "end_status" : "dropoff"  
  "ride_duration_in_secs" : 367  
  "reason" : "DROPOFF_SEEN"  
  "n_points" : 12  
}
```

- Have a Python development environment ready, with **Python 3.10**
- Use your favourite IDE (Visual Studio Code, PyCharm, ...)
- Clone the workshop repo

github.com/iht/beam-state-timers-quickstart

- Install dependencies with

```
pip install -r requirements.txt
```

- Run the tests (they fail until you write the solution 😊) with `pytest`.
- Have a look at the file `my_pipeline/pipeline.py` and fill the gaps



QUESTIONS?

Israel Herraiz
[linkedin.com/in/herraiz](https://www.linkedin.com/in/herraiz)
github.com/iht

Miren Esnaola
[linkedin.com/in/mirenesnaola](https://www.linkedin.com/in/mirenesnaola)
github.com/apichick