

# Exabyte-scale Streaming Iceberg IO with Beam, Ray, and DeltaCAT



- A Brief History of Amazon BI
- Apache Iceberg and DeltaCAT Overview
- The Iceberg Streaming Problem
- An Open Exabyte-Scale Solution
- Current State & Future Work

# A Brief History of Amazon BI

## 2016–2018 PB-Scale Oracle Data Warehouse Deprecation

- Migrated 50PB from Oracle Data Warehouse to S3-Based Data Catalog
- Decoupled storage with Amazon Redshift & Apache Hive on Amazon EMR Compute

## RIP Oracle Data Warehouse



← Post



Werner Vogels ✅  
@Werner

...

Amazon's Oracle data warehouse was one of the largest (if not THE largest) in the world. RIP. We have moved on to newer, faster, more reliable, more agile, more versatile technology at more lower cost and higher scale. #AWS Redshift FTW!



Andy Jassy ✅ 📧 @ajassy · Nov 9, 2018

In latest episode of "uh huh, keep talkin' Larry," Amazon's Consumer business turned off its Oracle data warehouse Nov 1 and moved to Redshift. By end of 2018, they'll have 88% of their Oracle DBs (and 97% of critical system DBs) moved to Aurora and DynamoDB. #DBFreedom

4:44 PM · Nov 9, 2018

625 Reposts 65 Quotes 1,524 Likes 31 Bookmarks

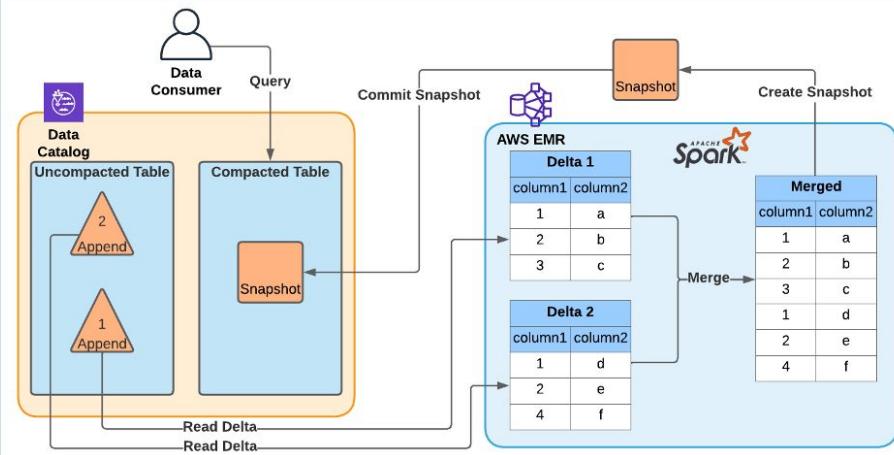


# A Brief History of Amazon BI

## 2018–2019 EB-Scale Data Catalog & Lakehouse Formation

- Bring your own compute (EMR Spark, AWS Glue, Amazon Redshift Spectrum, etc.)
- LSM-based CDC “Compaction” using Apache Spark on Amazon EMR

# Append-Only Compaction



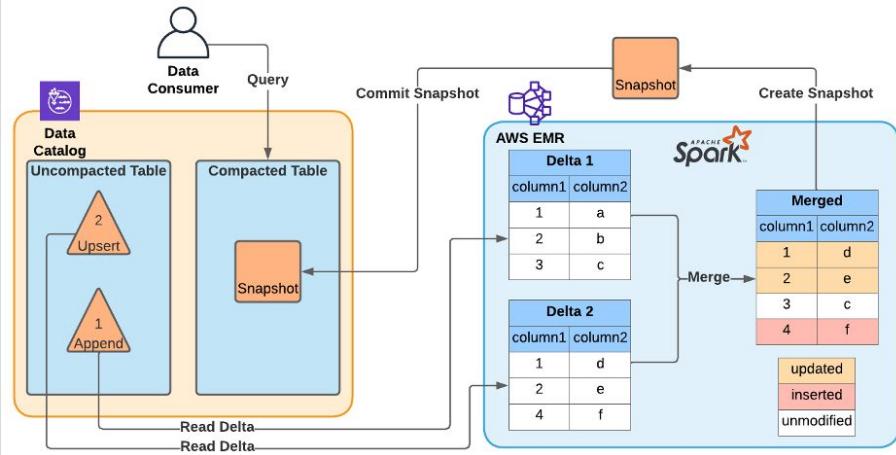
*Append deltas arrive in a table's CDC log stream, where each delta contains pointers to one or more S3 files containing records to insert into the table. During a compaction job, no records are updated or deleted so the delta merge is a simple concatenation, but the compactor is still responsible for writing out files sized appropriately to optimize reads (i.e. merge tiny files into larger files and split massive files into smaller files).*

# A Brief History of Amazon BI

## 2018–2019 EB-Scale Data Catalog & Lakehouse Formation

- Bring your own compute (EMR Spark, AWS Glue, Amazon Redshift Spectrum, etc.)
- LSM-based CDC “Compaction” using Apache Spark on Amazon EMR

# Upsert Compaction



*Append and Upsert deltas arrive in a table's CDC log stream, where each Upsert delta contains records to update or insert according to one or more merge keys. In this case, column1 is used as the merge key, so only the latest column2 updates are kept per distinct column1 value.*



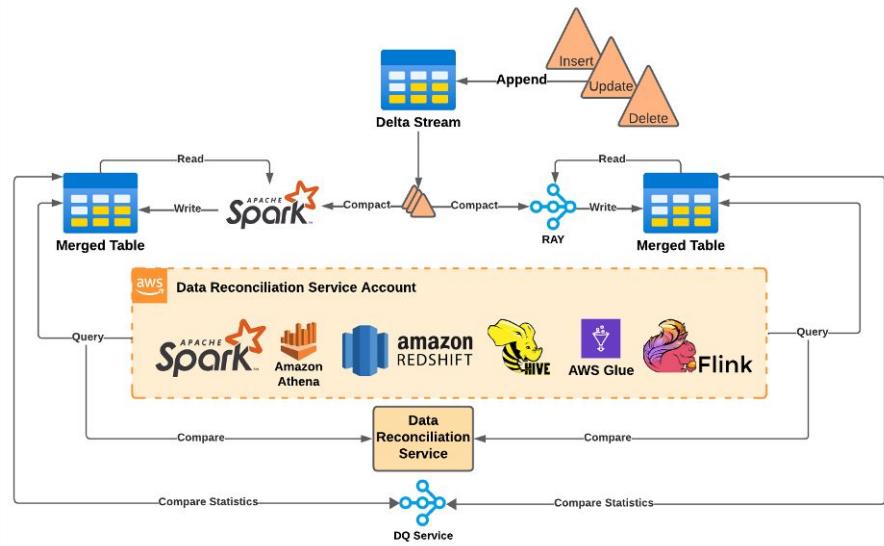
- **Pythonic**
  - Provides distributed Python APIs for ML, data science, and general workloads.
- **Intuitive**
  - Relatively simple to convert single-process Python to distributed.
- **Scalable**
  - Can integrate PB-scale datasets with data processing and ML pipelines.
- **Performant**
  - Reduces end-to-end latency of data processing and ML workflows.
- **Efficient**
  - Reduces end-to-end cost of data processing and ML.
- **Unified**
  - Can run all steps of mixed data processing, data science, and ML pipelines.

# A Brief History of Amazon BI

## 2019–2023 Ray Integration

- EB-Scale Data Quality Analysis
- Spark-to-Ray Compaction Migration
- Reduced Cost by 82% (equivalent to \$120MM/year of Amazon EC2 on-demand charges)

# Ray Shadow Compaction



New deltas arriving in a data catalog table's CDC log stream are merged into two separate compacted tables maintained separately by Apache Spark and Ray. The Data Reconciliation Service verifies that different data processing frameworks produce equivalent results when querying datasets produced by Apache Spark and Ray, while the Ray-based DQ Service compares key dataset statistics.

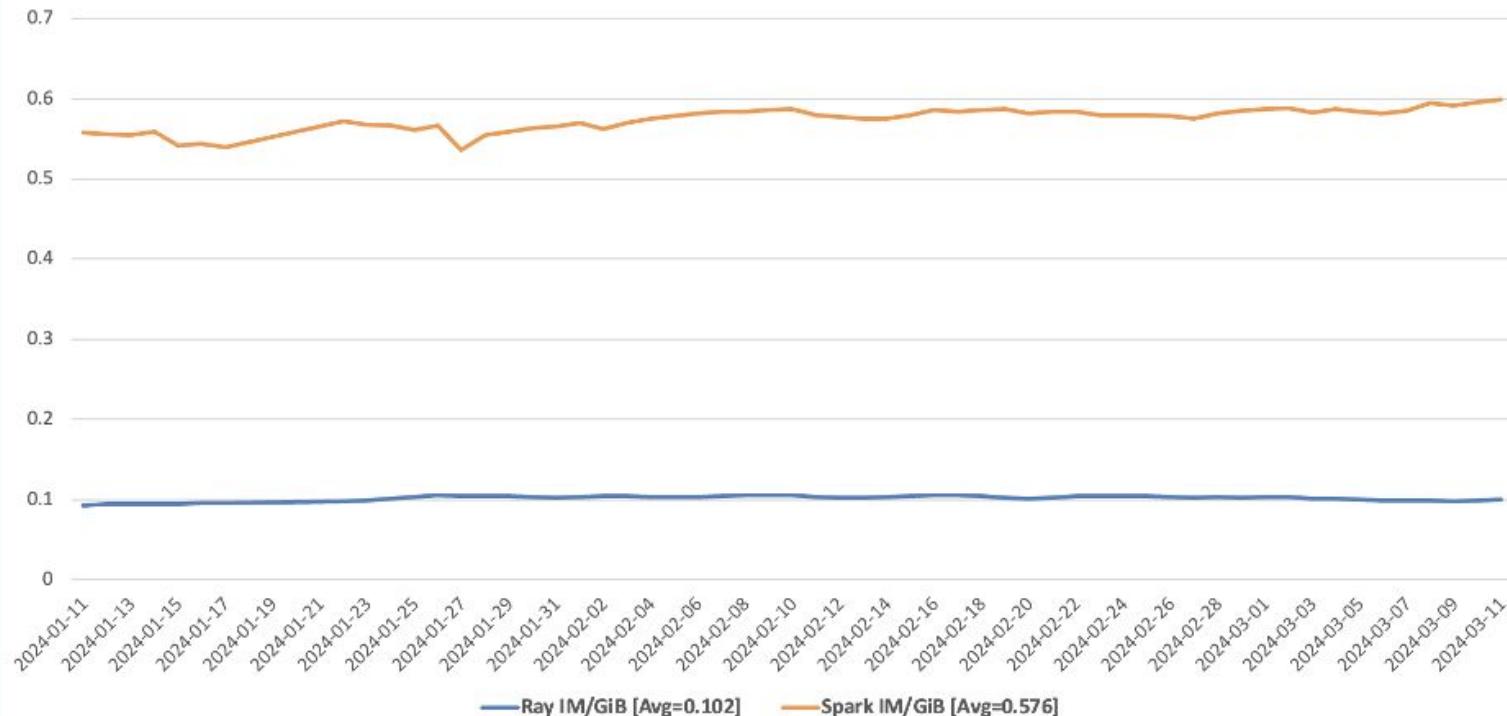


# Ray vs. Spark Compactor Efficiency

Ray vs. Spark Efficiency: EC2 Instance Minutes (IM) Per GiB Compacted

>1.2EiB of Input S3 Parquet Data

2024-01-11 to 2024-03-11

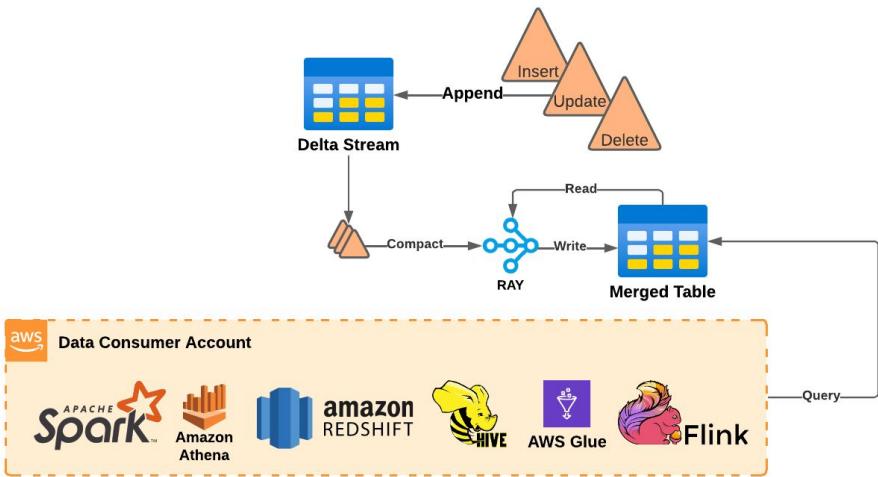


## A Brief History of Amazon BI

# 2024-2025 Ray Exclusivity

- Migrate all Table Queries to Ray Compactor Output
- Turn off Spark Compactor
- OSS implementation of Ray Compactor in DeltaCAT

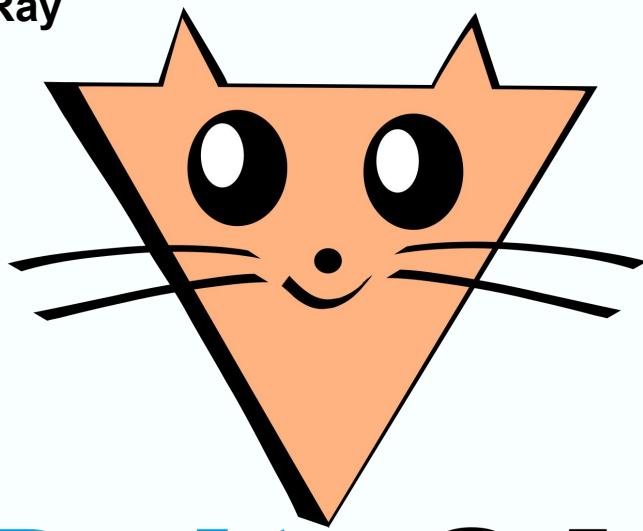
## Ray Exclusive Compaction



*New deltas arriving in a data catalog table's CDC log stream are merged into only one compacted table maintained by Ray. Amazon BI tables are gradually being migrated from Spark compaction to Ray-Exclusive compaction, starting with our largest tables.*

## A Portable Pythonic Data Lakehouse Powered by Ray

- **Catalog:** High-level APIs to create, discover, organize, share, and manage datasets.
- **Compute:** Distributed data management procedures to read, write, and optimize datasets.
- **Storage:** In-memory and on-disk multimodal dataset formats.
- **Sync:** Synchronize DeltaCAT datasets to data warehouses and other table formats.



v1.X Used in Production @ Amazon (But Also Overfit to Amazon's Use-Case)

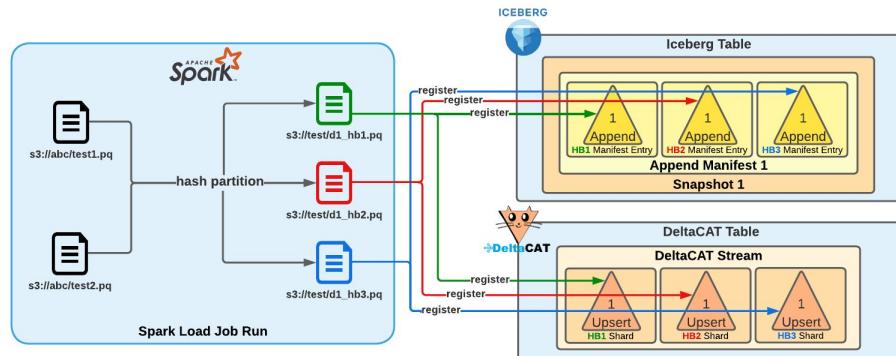
v2.0 in Development for General Purpose Use (e.g. Iceberg Table Management)

# A Brief History of Amazon BI

## 2025+ Lakehouse Unification

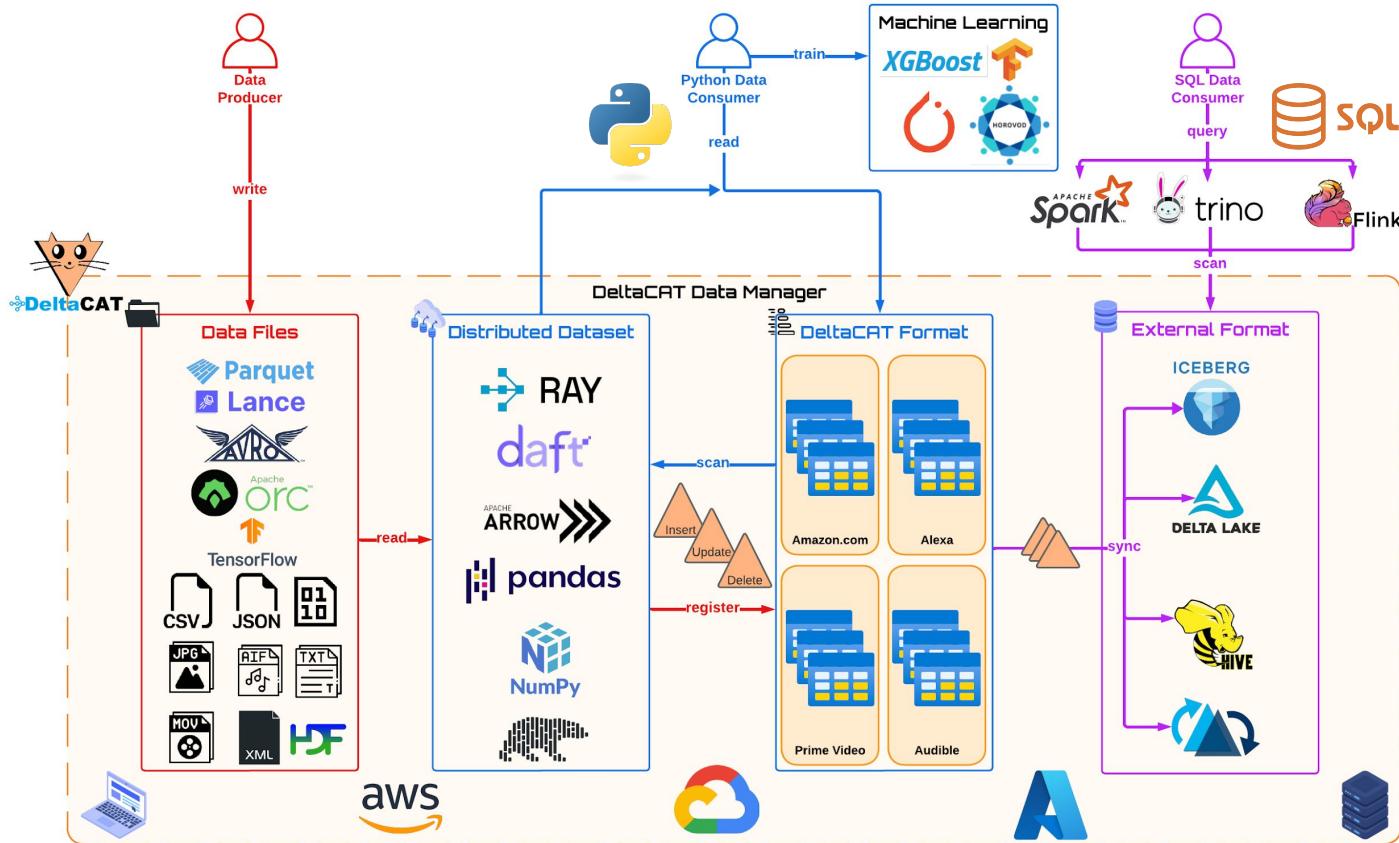
- Reuse data files across multiple table formats
- Batch & Streaming Compute Integration with Iceberg
- OSS Implementations of Iceberg Table Maintenance Jobs using Ray in DeltaCAT

# Shared Data Files

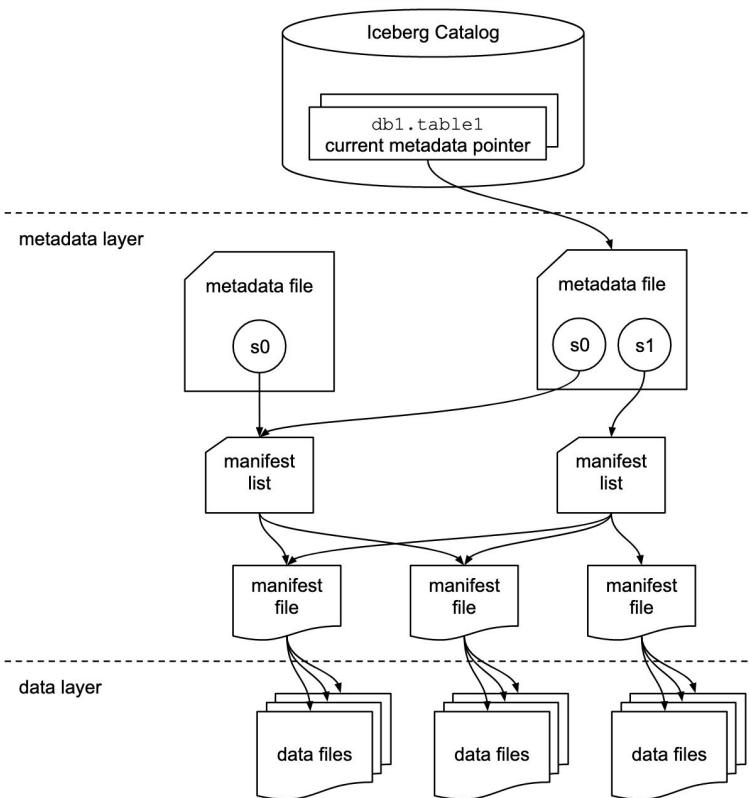


*New deltas arriving in a data catalog table's CDC log stream reuse the same data files across Iceberg, DeltaCAT, and any other compatible table formats (e.g., Hudi, Hive, etc.).*

# DeltaCAT Overview



# Iceberg Overview



- **Table Metadata Format**

- Table Metadata File
- Snapshot File
- Manifest File
- Data File

- **Catalog**

- Table Metadata Pointer
- Java & Python **Interfaces**
- REST/Hive/JDBC/Glue/etc. **Implementations**

# Iceberg Table Directory

d\_mp\_asins\_na\_bucket\_asin\_1000/

Objects Properties

## Objects (2) [Info](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to

Find objects by prefix

<input type="checkbox"/>	Name	Type
<input type="checkbox"/>	data/	Folder
<input type="checkbox"/>	metadata/	Folder

# Iceberg Metadata Directory

metadata/

Objects Properties

## Objects (5) [Info](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to

Find objects by prefix

<input type="checkbox"/>	Name	Type
<input type="checkbox"/>	00000-69d58ad1-0b40-4b2d-8c70-e718c3f2eaea.metadata.json	json
<input type="checkbox"/>	00001-8ff5a00c-bd3e-44c3-8345-47d633e7a368.metadata.json	json
<input type="checkbox"/>	00002-aed671ae-db3d-4c5a-aefb-4801b2b4dbdd.metadata.json	json
<input type="checkbox"/>	a3d06987-8f93-42d4-96d2-ad1aee619474-m0.avro	avro
<input type="checkbox"/>	snap-3681555070106584832-1-a3d06987-8f93-42d4-96d2-ad1aee619474.avro	avro

# Iceberg Data Directory

data/

**Objects** **Properties**

To enable sorting in the table below, use the search to reduce the size of the list

**Objects (999+)** [Info](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#)

[Find objects by prefix](#)

<input type="checkbox"/>	Name	Type
<input type="checkbox"/>	asin_bucket_1000=0/	Folder
<input type="checkbox"/>	asin_bucket_1000=1/	Folder
<input type="checkbox"/>	asin_bucket_1000=10/	Folder
<input type="checkbox"/>	asin_bucket_1000=100/	Folder
<input type="checkbox"/>	asin_bucket_1000=101/	Folder
<input type="checkbox"/>	asin_bucket_1000=102/	Folder

# Iceberg Partition Directory

asin\_bucket\_1000=10/

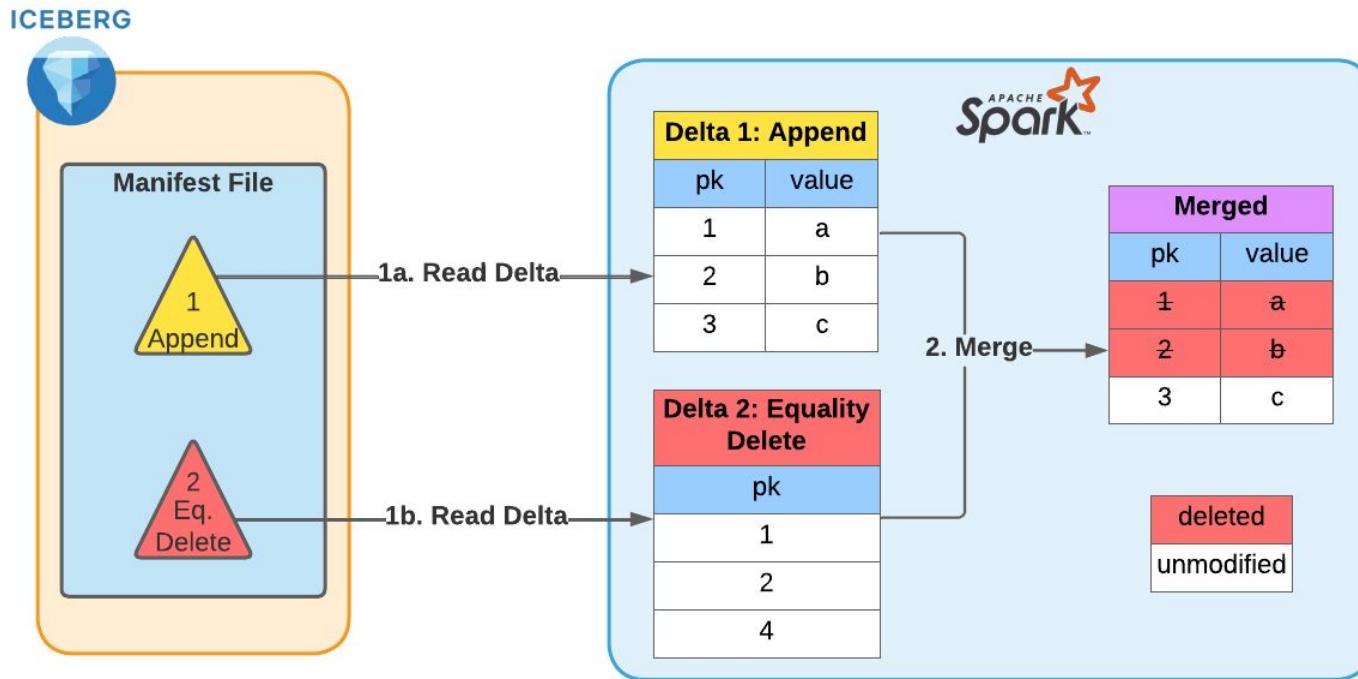
**Objects (11)** [Info](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a

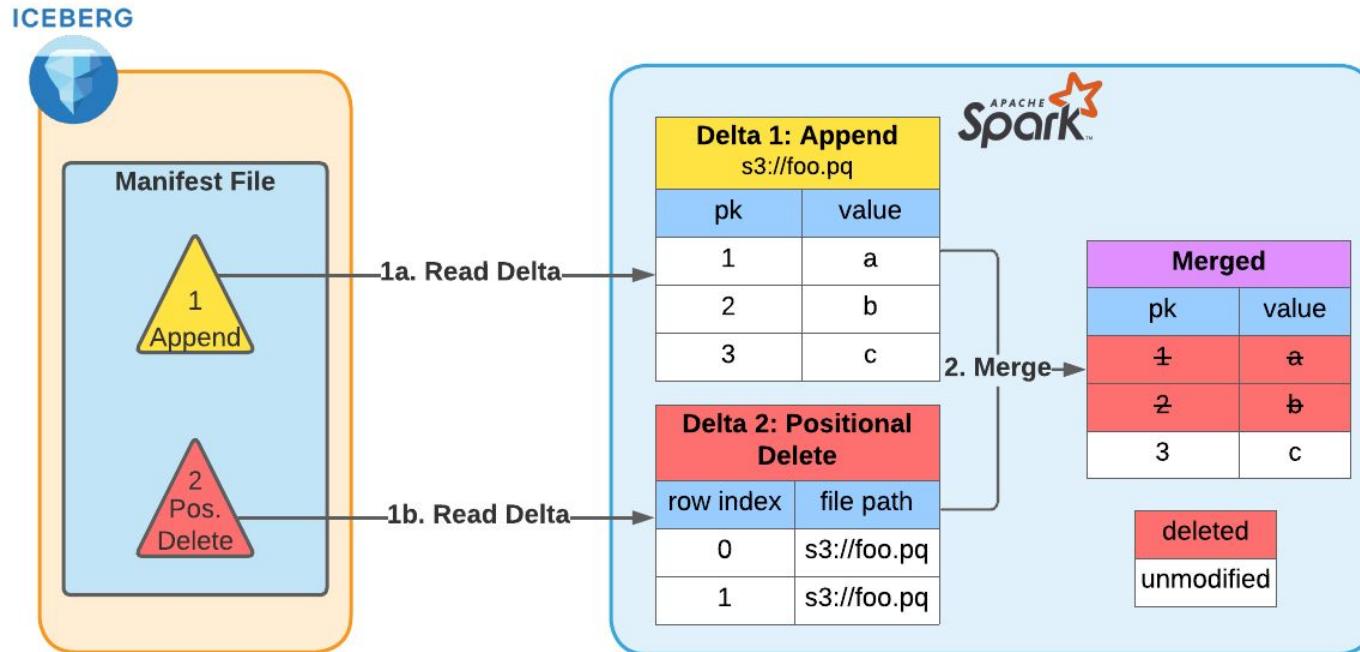
[Find objects by prefix](#)

<input type="checkbox"/>	Name	Type
<input type="checkbox"/>	00226-63726-9f3d1f7d-20f6-4a2f-b325-b6f156a780c1-00001.parquet	parquet
<input type="checkbox"/>	00226-63726-9f3d1f7d-20f6-4a2f-b325-b6f156a780c1-00002.parquet	parquet
<input type="checkbox"/>	00226-63726-9f3d1f7d-20f6-4a2f-b325-b6f156a780c1-00003.parquet	parquet
<input type="checkbox"/>	00226-63726-9f3d1f7d-20f6-4a2f-b325-b6f156a780c1-00004.parquet	parquet
<input type="checkbox"/>	00226-63726-9f3d1f7d-20f6-4a2f-b325-b6f156a780c1-00005.parquet	parquet

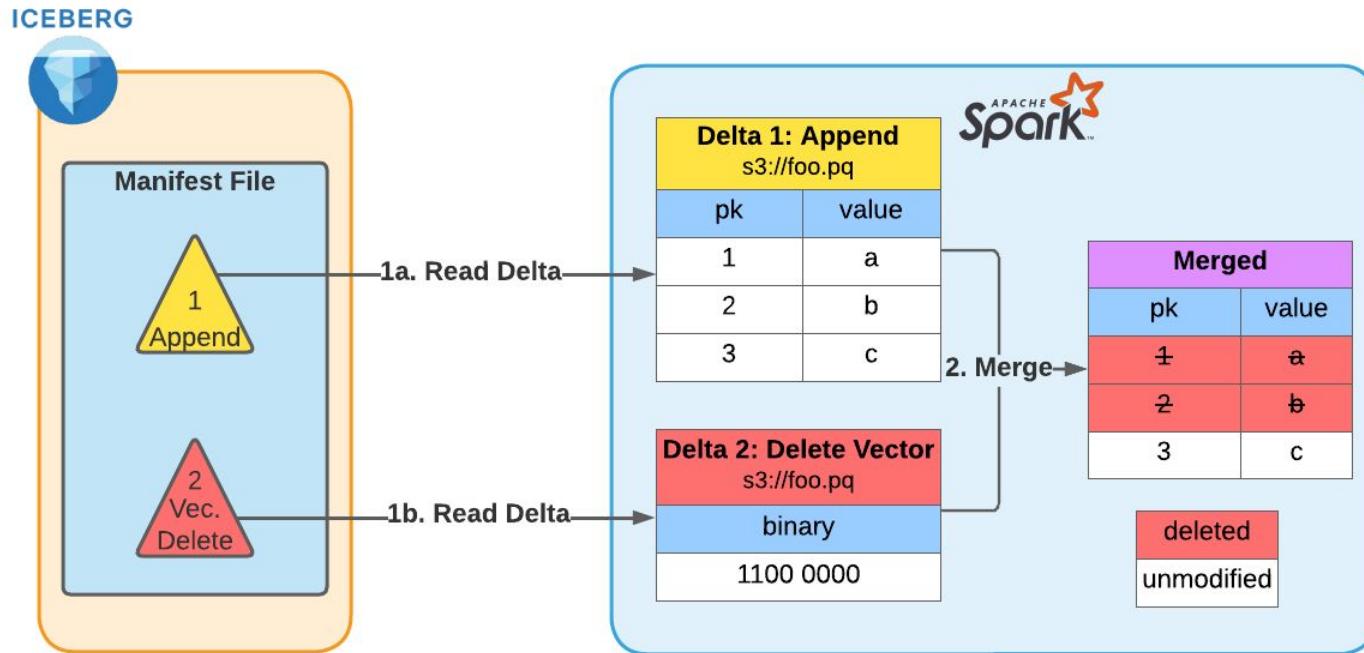
# Iceberg Equality Deletes



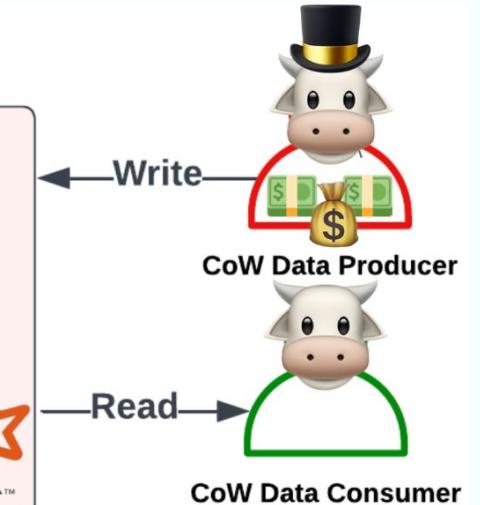
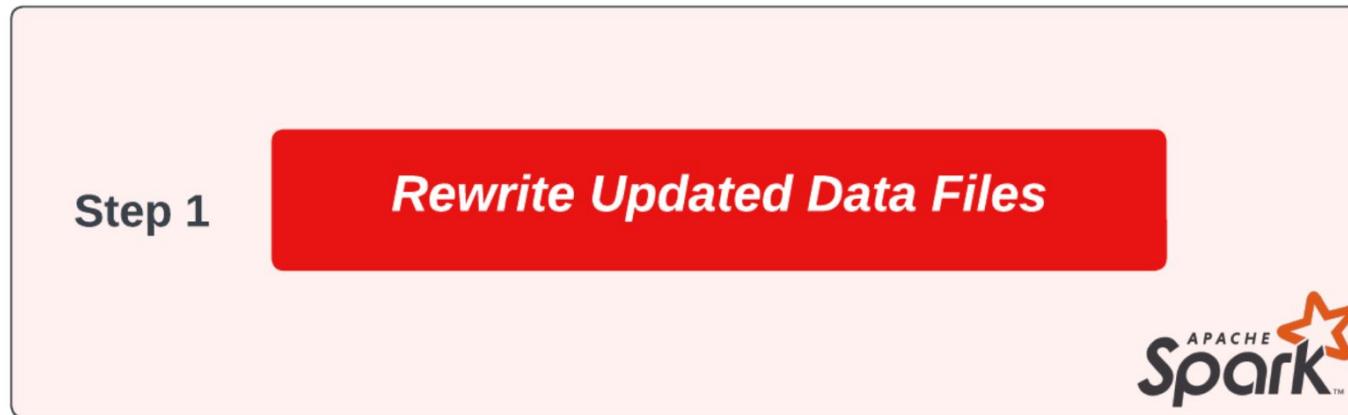
# Iceberg Positional Deletes



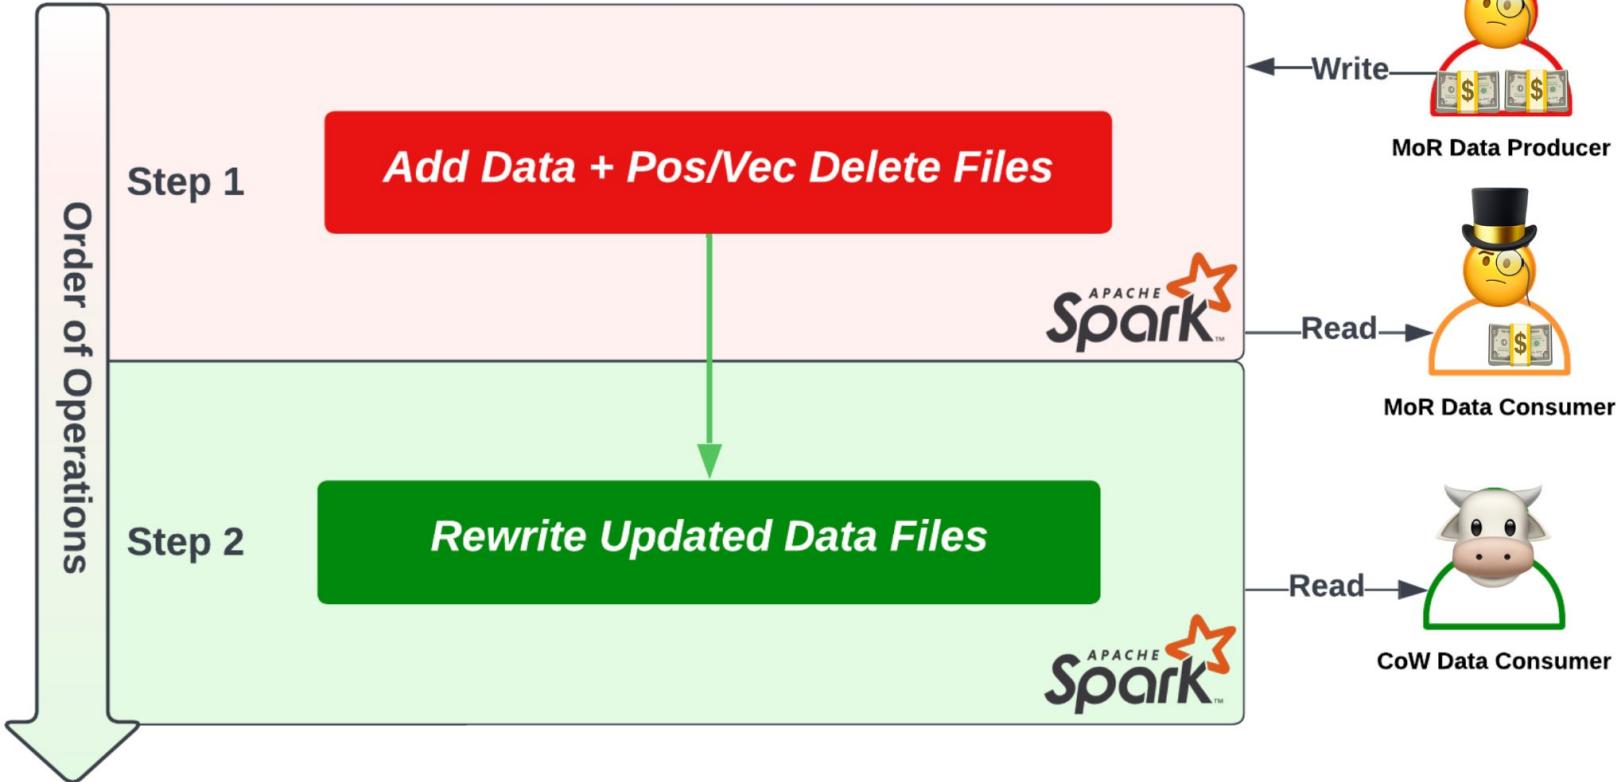
# Iceberg Delete Vectors



# Iceberg Copy on Write



# Iceberg Merge on Read



# The Iceberg Streaming Problem

- Merge-on-Read Tables
  - **Streaming** Frameworks like Flink Prefer to Write
    - **Equality Deletes** (Cheap Writes, Expensive Reads)
    - **Fast & Conflict-Free** but may cause **OOM Errors** on Read 😞
  - **Batch** Frameworks like Spark Write
    - **Positional Deletes** (IcebergV2)
    - **Binary Delete Vectors** (IcebergV3)
    - Moderately Expensive Writes, Cheaper Reads
    - **Less practical for High-Frequency Writes** at TB-PB Scale 😞
    - Why? **High-Latency** and Susceptible to **Irresolvable Write Conflicts**
- Copy-on-Write Tables
  - **Always Rewrite Data Files** w/ Deletes Applied (Cheapest Reads, Expensive Writes)
  - **Impractical for High-Frequency Writes** at TB-PB Scale 😞

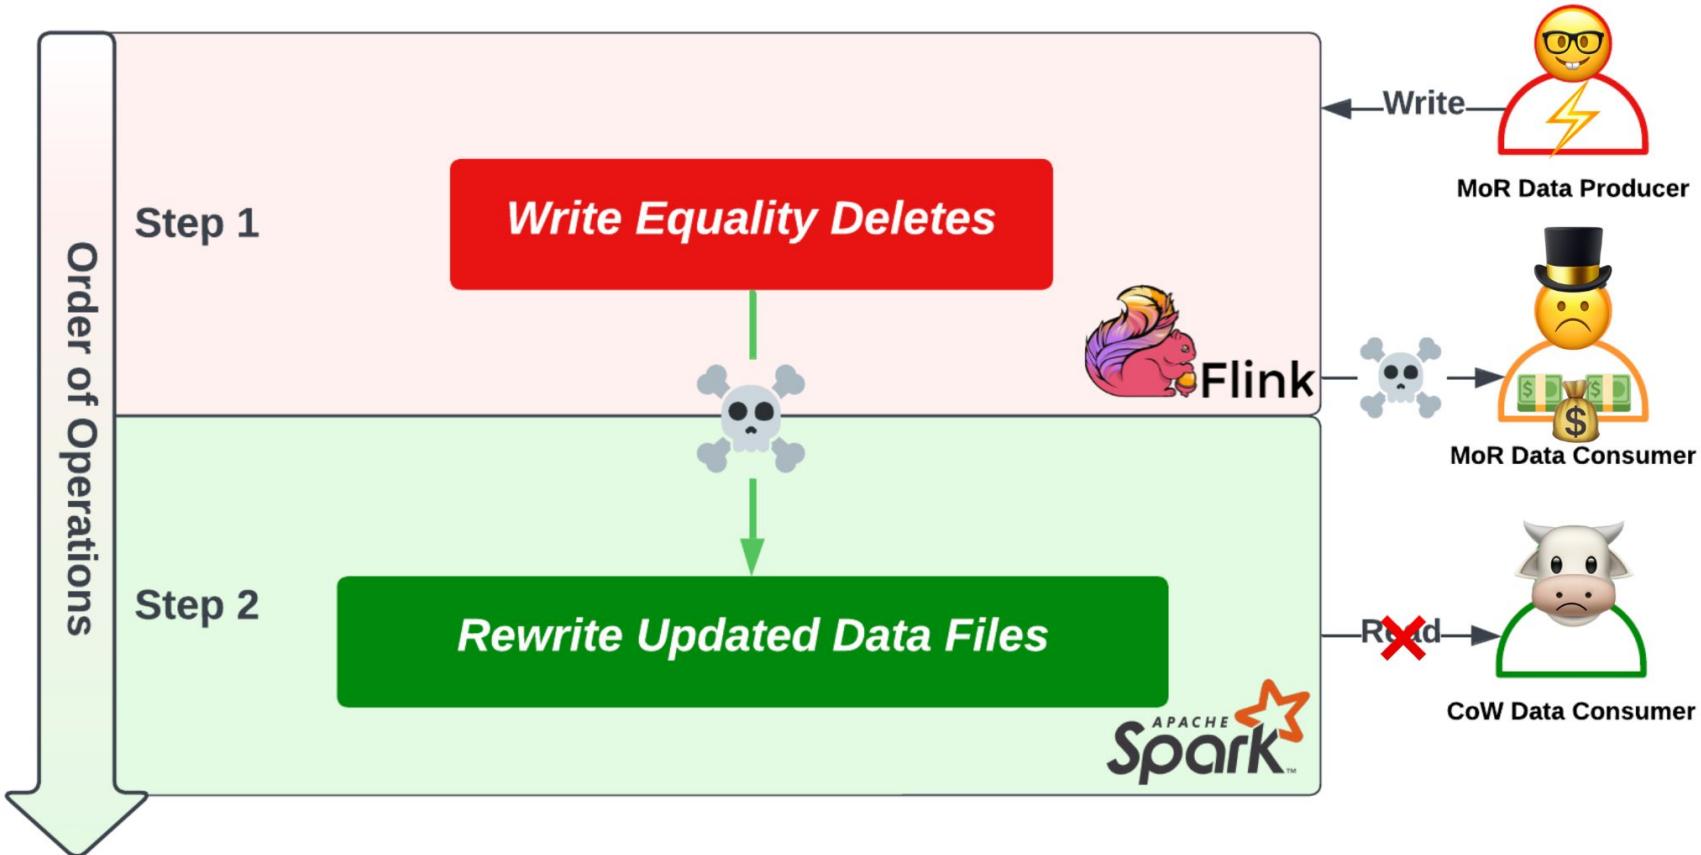


# The Iceberg Streaming Problem

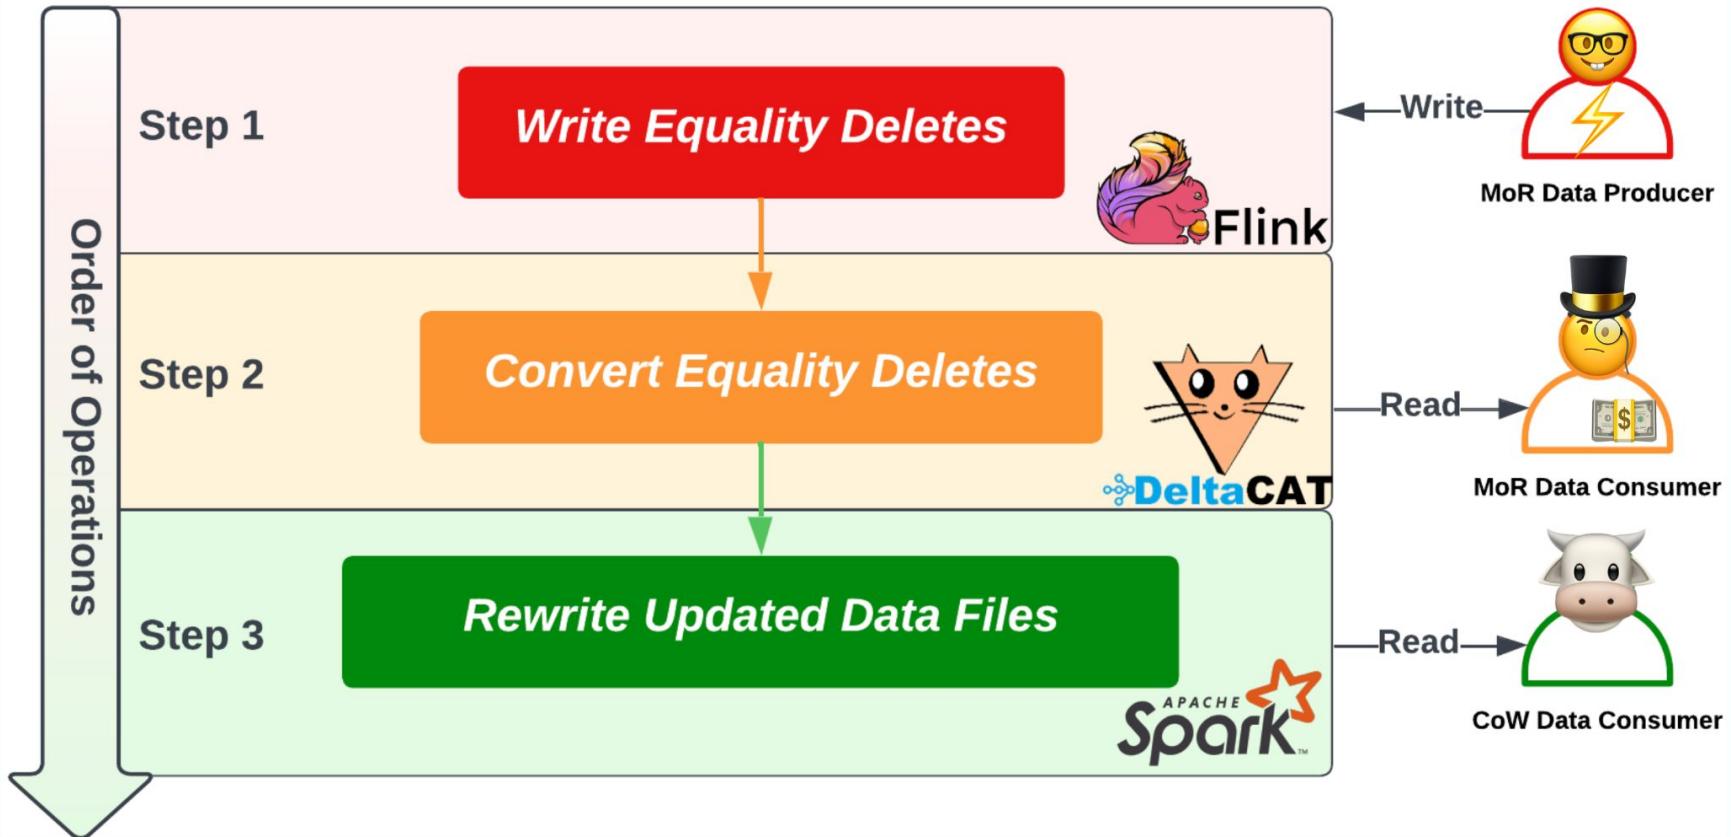
- **Issues & PRs Raised and Abandoned Since 2020**

- 2020-05-08: [Add an action to rewrite equality deletes as position deletes](#)
- 2021-02-04: [Spark: support replace equality deletes to position deletes](#)
- 2021-03-23: [Add an action to rewrite equality deletes](#)
- 2023-02-27: [Data file rewriting spark job fails with oom](#)
- 2023-12-04: [RewritePositionDeleteFiles cannot work with equality delete file?](#)
- 2024-03-27: [Spark rewrite Files Action OOM](#)
- 2024-07-09: [Add RocksDBStructLikeSet for storing equality deletes](#)

# The Iceberg Streaming Problem



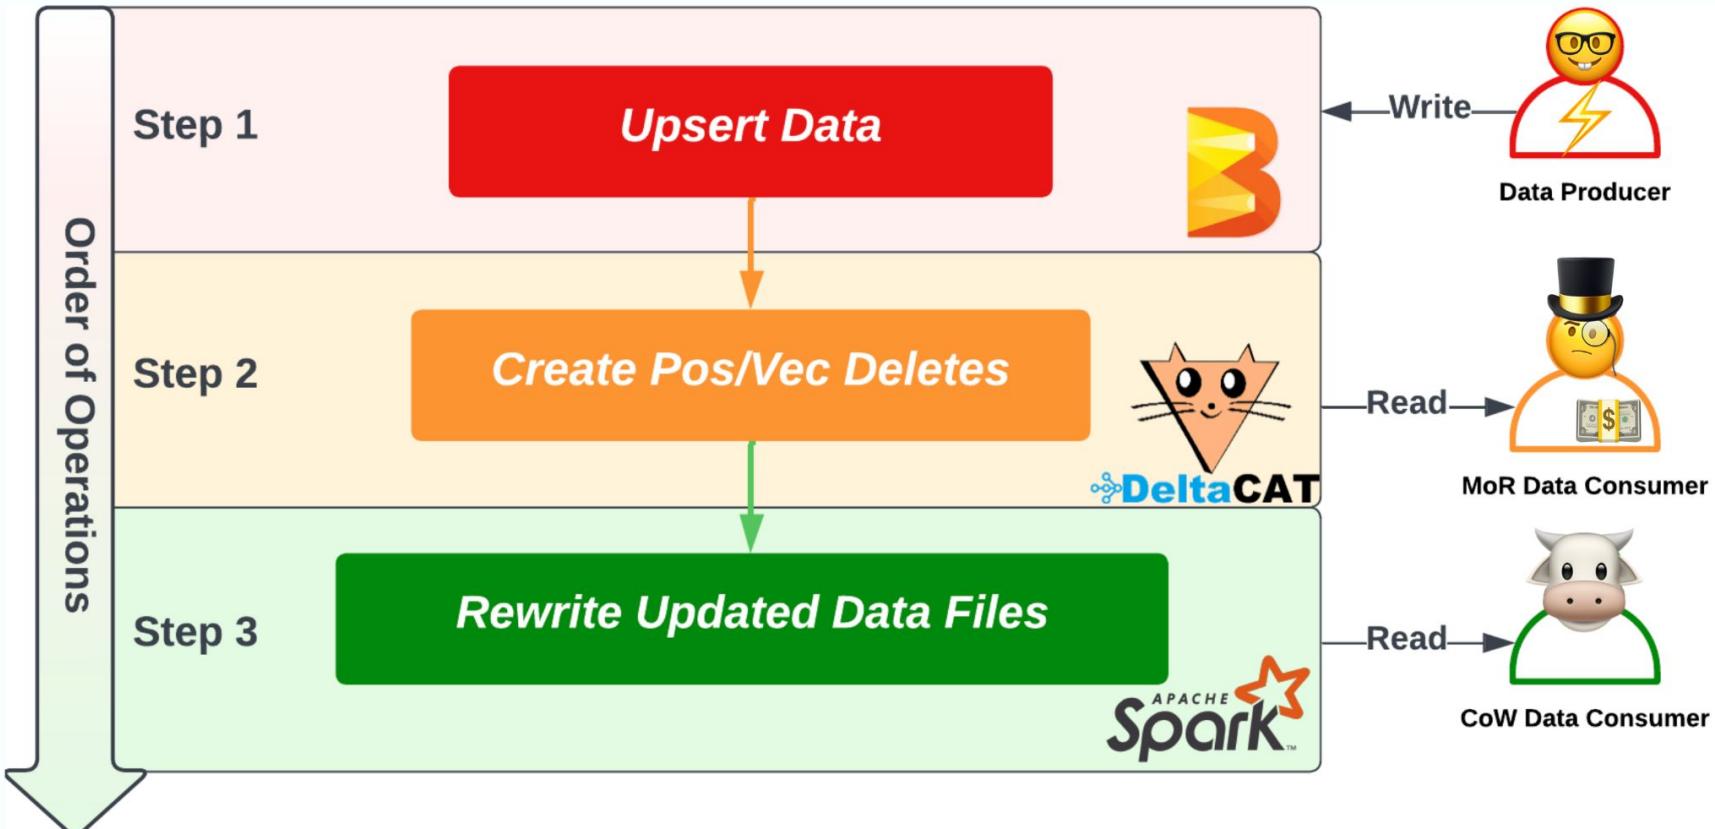
# An Open Exabyte-Scale Solution



# An Open Exabyte-Scale Solution



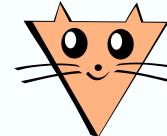
# An Open Exabyte-Scale Solution





## Testing with PB-scale production Iceberg tables

**Q4 2025:** Production onboarding to **EB-scale** production data catalog



ΔeltaCAT

## Python converter job ready for experimental use

**Beam** Iceberg **upserts** via `beam.managed.Read() /Write()` wrapper

**Flink** Iceberg **delete conversion** (equality to positional/vector)

Run **manually** or **automatically** via Iceberg table monitor agent

Run **locally** or on a **distributed** Ray cluster (in GCP, AWS, etc.)



<https://github.com/ray-project/deltacat/tree/2.0/deltacat/examples/experimental/iceberg/converter/>

- **Python Native Beam IO Connector?**
  - +1 Github Issue #561 @ <https://github.com/ray-project/deltacat/issues/561>
- **Flink IO Connector?**
  - +1 Github issue #562 @ <https://github.com/ray-project/deltacat/issues/562>
- **Ray Positional/Vector Delete Materialization for Iceberg?**
  - +1 Github issue #121 @ <https://github.com/ray-project/deltacat/issues/121>

Patrick Ames

# QUESTIONS?

@Patrick Ames on Ray Slack

<https://github.com/ray-project/deltacat>

<https://github.com/pdames>