

Optimize parallelism for reading from  
Apache Kafka to Dataflow

## Introduction

Consumption from Kafka suddenly stopped!

Messages from Kafka are being lost/dropped in the Dataflow pipeline


Most of the workers show very little CPU utilization



Supriya Koppa

*Technical Support Specialist,  
Google Cloud*

## WHAT GIVES PEOPLE FEELINGS OF POWER






# Agenda



- Introduction
- Kafka
  - Fundamentals
  - Parallelism in Kafka
- Dataflow
  - Fundamentals
  - Dataflow as a Kafka consumer
- The challenge of parallelism
  - Common complaints and attempted mitigations
  - And their symptoms
- Recommended ways to unlock parallelism
- Let's review
- Questions... Thoughts...


# Kafka fundamentals

- > Producers
- > Kafka cluster (broker, topic)
- > Kafka Topic (partition, replication)
- > Consumers (consumer group, consumer offset)




# Parallelism in Kafka

- > Parallelism is inherently tied to Partitions
- > Generally, more the partitions, higher the processing parallelism




# Dataflow fundamentals

Apache Beam SDK  
Managed  
Scalable  
Stream and Batch  
Exactly-once processing  
GCP products  
Monitoring



Storage  
Analysis  
Visualization



# Dataflow as a Kafka consumer

## > Managed I/O

```
#Code sample only for demonstrative use
PCollection<String> kafkaRecords = pipeline.apply(
    "ReadFromKafkaManaged",
    ManagedIO.read(
        ManagedIO.SourceSpec.of("KAFKA") // Specify the source type as KAFKA
        .withConfiguration(ImmutableMap.of(
            "topic", "your-kafka-topic-managed",
            "bootstrap.servers", "your-kafka-broker:9092",
            "consumer.group.id", "my-dataflow-managed-group"
        // Dataflow manages many other Kafka-specific configurations internally
        ))
    )
);
```

## > KafkaIO

```
#Code sample only for demonstrative use
PCollection<KafkaRecord<String, String>> kafkaRecords = pipeline.apply(
    "ReadFromKafkaIO",
    KafkaIO.<String, String>.read()
        .withBootstrapServers("your-kafka-broker:9092")
        .withTopics(Collections.singletonList("your-kafka-topic-io"))
        .withConsumerConfigUpdates(ImmutableMap.of(
            "group.id", "my-dataflow-kafkaio-group",
            "max.poll.records", "5000",           // Explicitly tune batch size
            "fetch.max.bytes", "10485760"        // Explicitly tune fetch size (10MB)
        // Add other Kafka consumer properties as needed for tuning
        ))
    // Optional: Enable redistribution for parallelism beyond partition count
    .withRedistribute()
    .withKeyDeserializer(StringDeserializer.class)
    .withValueDeserializer(StringDeserializer.class)
    .withoutMetadata() // If you only need key/value, not offset/timestamp
);
```

## Challenge of parallelism

“Reading from Kafka is slow, or has stopped”

“CPU usage of Dataflow workers is skewed”

“The number of partitions was increased, but it did not resolve the problem”

## Mitigations commonly attempted

“Restarting the job”

“Adding more workers, or better workers”

# What the symptoms look like...

Google Cloud kafka-dataflow-BeamSummit2025

Search (/) for resources, docs, products, and more

Job info

JOB GRAPH EXECUTION DETAILS JOB METRICS COST RECOMMENDATIONS AUTOSCALING

Job steps view  
Graph view

CLEAR SELECTION

Job name: kafka-bigquery  
Job ID: 2025-06-17\_23\_59\_40-15177191742749806522  
Job type: Streaming  
Job status: Running  
SDK version: Apache Beam SDK for Java 2.65.0  
Job region: us-central1  
Service zones: us-central1-f  
Worker location: us-central1  
Current workers: 1  
Latest worker status: Worker pool started.  
Straggler status: No active straggler  
Start time: June 17, 2025 at 12:29:41 PM GMT+5  
Elapsed time: 7 days 8 hr  
Encryption type: Google-managed  
Dataflow Prime: Disabled  
Dataplex Lineage: Disabled  
Runner v2: Disabled  
Streaming Engine: Enabled  
Vertical Autoscaling: Disabled  
Streaming Mode: Exactly once

Show debug panel

Logs HIDE

JOB LOGS WORKER LOGS DIAGNOSTICS DATA SAMPLING

Severity: Info

Filter: Search all fields and values

1 HOUR

Resource metrics

Current vCPUs: 2  
Total vCPU time: 352.789 vCPU hr  
Current memory: 7.5 GB  
Total memory time: 1,322.959 GB hr  
Current HDD PD: 30 GB  
Total HDD PD time: 5,291.836 GB hr

Release Notes

# BEAM SUMMIT NYC 2025

# Recommended ways to unlock parallelism

## Kafka

### > Partitions

*Golden rule*

```
partitions = 2 * number_of_vCPUs
```

Example: n1-standard-4, 10 max workers.

partitions =  $2 * 4 * 10 = 80$

```
dataflow_fanout = 4 * max_num_workers
```

partitions =  $4 * 10 = 40$

### > Hot partition




# Recommended ways to unlock parallelism


## Dataflow

> Add redistribute

Replace `KafkaIO.Read`




with `KafkaIO.Read.withRedistribute`



> `KafkaIO.Read.withRedistributeNumKeys`

> Autoscaling (monitoring metrics)



## Recommended ways to unlock parallelism

### KafkaIO

```
> unboundedReaderMaxElements  
> unboundedReaderMaxReadTimeMs  
In DataflowPipelineDebugOptions
```

```
> max.poll.records  
> fetch.max.bytes  
> max.partition.fetch.bytes  
In KafkaIO.Read<K,V> withConsumerConfigUpdates
```

```
> consumerPollingTimeout  
In KafkaIO.Read<K,V>  
withConsumerPollingTimeout
```

# Let's review

## Kafka

- > Partitions

*Golden rule*

```
partitions = 2 * number_of_vCPUs
```

Example: n1-standard-4, 10 max workers.  
partitions =  $2 * 4 * 10 = 80$

```
dataflow_fanout = 4 * max_num_workers
```

partitions =  $4 * 10 = 40$

- > Hot partition

## Dataflow

- > Add redistribute

Replace `KafkaIO.Read`

with `KafkaIO.Read.withRedistribute`

```
KafkaIO.Read.withRedistributeNumKeys
```

- > Autoscaling (monitoring metrics)

## KafkaIO

```
> unboundedReaderMaxElements  
> unboundedReaderMaxReadTimeMs  
In DataflowPipelineDebugOptions
```

```
> max.poll.records  
> fetch.max.bytes  
> max.partition.fetch.bytes  
In KafkaIO.Read<K,V>  
withConsumerConfigUpdates
```

```
> consumerPollingTimeout  
In KafkaIO.Read<K,V>  
withConsumerPollingTimeout
```

Thank you for listening!

QUESTIONS?  
Thoughts?